

# Conceptual Loops for Enterprise Document AI


Reduce uncertainty. Don't waste experience.

Dr. Jordy Van Landeghem · Probably Approximately Human

# whoami

- Research intern **ORACLE** (Seq2Seq for Dialogue Modeling) and **NUANCE** (LM Algorithms)
- Founding AI Research Engineer **contract.fit** [2017-2024]
- PhD project: *Intelligent Automation for AI-driven Document Understanding* **KU LEUVEN**
- Model Dev - GenAI - Agents Technical Lead **INSTABASE** [2024-...]
- Founded my own company  **Probably Approximately Human** [2026]

 Belgium 



# You might (hopefully) know me from ...

## Document UnderstanDing of Everything (**DUDE**) 🤖

### #non-answerable

Q: In which year does the Net Requirement exceed 25,000?

A: None

### #abstractive #counting

Q: How many attorneys are listed for the plaintiffs?

A: Two

### #layout-navigating #graphic-intensive

Q: Are the margins of the page uniform on all pages?

A: Yes

### #extractive #list

Q: What are the Years mentioned in Chart 1?

A: [2020, 2021, 2022]



Page 1



Page 2



Page N

### #multi-hop #layout-navigating

Q: From the list of Top 10 Key Recovery Components, which is the last component listed on the second page?

A: Hope


### #abstractive #graphic-intensive

Q: Does this document contain any checkboxes?

A: No

Or more recently....

***MADQA / Strategic Navigation or Stochastic Search? How Agents and Humans Reason Over Document Collections***



# Agenda

---



## 01

### Reality checks from industry

From benchmark performance to production automation

Why  $f(x) = y$  stops being enough, and why a loop is needed at all

## 02

### The inner loop

1. Representation
2. Search
3. Evidence
4. Action

*Where agents can help, one uncertainty question per stage*

## 03

### The outer loop

- Experience & Feedback
- Memory (remember, promote, compile)
- Don't waste experience
- Open Research frontier



## Section 01

# Reality checks

Traditional DocAI vs. Modern Requirements

01

# Two very different objectives

## DOCUMENT UNDERSTANDING

$$f(x) \rightarrow y$$

|                 |                       |
|-----------------|-----------------------|
| <b>Goal</b>     | benchmark performance |
| <b>Unit</b>     | model / prediction    |
| <b>Accuracy</b> | primary metric        |
| <b>Failure</b>  | wrong prediction      |
| <b>Logic</b>    | mostly learned        |
| <b>Learning</b> | train → deploy        |

## ENTERPRISE DOCUMENT AI

document → evidence → action → outcomes

|                 |                                                    |
|-----------------|----------------------------------------------------|
| <b>Goal</b>     | automation + trust                                 |
| <b>Unit</b>     | end-to-end system                                  |
| <b>Accuracy</b> | upper bound on automation                          |
| <b>Failure</b>  | wrong + what happens next? (verify/trace/escalate) |
| <b>Logic</b>    | prediction + rules + workflow                      |
| <b>Learning</b> | experience → feedback → improve                    |

A benchmark asks, “Was the answer right?”

Production asks, “Can the process run reliably, cheaply, auditably, and at scale?”

# What form does $f(x)$ typically take?

## VLMs

- extract via a single model call
- Closed API models & open-weight models

```
from google.genai import types
with open(file_path, "rb") as f:
    data = f.read()
part = types.Part.from_bytes(data=data,
mime_type=mime_type)
```

`f(x)_per_field.json`

## Schema Output (Trust Values)

```
{
  "value": "Extracted Data",
  "confidence": 0.97,
  "provenance": {
    "source_page": 1,
    "bounding_box": [120, 45, 200, 80]
  }
}
```



GRITS

# 78.4 → 96.0

## after removing the continuous confidence request

**self-judging every value while transcribing a long table created an escape hatch.**

[illegible]

# New problems to solve, stochastic control over loops

## Enterprise Reliability Requirements

- **Automation** (adaptivity)
- **Defensible failure handling** (confidence/validations)
- **Auditability** (provenance)
- **Long context handling** (structured transcription)
- **Determinism** (stable behavior)
- **Continual learning** (production flywheel)

## The Paradigm Shift

**Reliable DocAI is built around stochastic intelligence.**

The future belongs to systems that spend intelligence where it counts, adapt strategy under uncertainty and explain why we should trust the result.

## For Researchers

### Measurable questions across the loops



This paradigm opens critical new paths to explore representation, retrieval, evidence, decisions, and learning throughout the execution cycle.

# The **agent loop**, inside the **feedback loop**

Before full autonomy — it's a reliable, engineered system.

## FEEDBACK LOOP — THE SYSTEM LEARNS

### AGENT LOOP — THE WORK HAPPENS





## Section 02


# Inner Loop

REPRESENTATION

What is here?

# What is actually in the document?

Can we represent the document faithfully?



<doclang/>

ParseBench

text

layout

tables

charts

formatting

reading order

visual relations


spatial semantics

## OBSERVATION

No method consistently dominates tables, charts, content faithfulness, formatting, and visual grounding.

# How would you read this complex document?

original



 upscal



en, ar



PaddleOCR

Google Cloud




 Gemini  Claude  ChatGPT




 crop

  denois

e



# So we built an Agent that *thinks with the documents*



```
Quality_assessment:
- Overall_quality: Medium
- Text_readability: Fair
- Image_orientation: Incorrect (rotated)
- Color_quality: Washed out
```

```
Preprocessing recommendations:
- Page Number 1:
  - Apply minor noise reduction to
    remove faint lines on the left
    margin (optional).
  - Rotate the image 12° clockwise.
  - Apply table extraction to capture
    transaction data.
  - Decode postal barcode if its
    information is relevant.
```

**Preprocessing Tools** are applied on each page wherever required on **parallel**.

**Not all pages or documents get processed, processing happens only where necessary**

Tools at our disposal:

- Denoise Tool
- OCR Engines
- Table Recognition
- Deskew Tool
- Specialized VLMs
- ...

**Research question:** Can agents optimize the accuracy / latency / cost frontier per document?



## Section 02



# Inner Loop

SEARCH

Where is it?



# Where is what I need?



Same fundamental operation. Different scale.

# Strategic Navigation or Stochastic Search?


## How Agents and Humans Reason Over Document Collections

Lukasz Borchmann ❄️ Jordy Van Landeghem 🌈 Michał Turski ❄️ Shreyansh Padarha 🏠  
Ryan Othniel Kearns 🏠 Adam Mahdi 🏠 Niels Rogge 🧐 Clémentine Fourrier 🧐  
Siwei Han 🏠 Huaxiu Yao 🏠 Artemis Llabrés <sup>CUC\*</sup> Yiming Xu <sup>CUC\*</sup> Dimosthenis Karatzas <sup>CUC\*</sup>  
Hao Zhang ❄️ Anupam Datta ❄️

🏆 Snowflake/MADQA-Leaderboard </> Snowflake-Labs/MADQA 🗄️ OxRML/MADQA

### Abstract

Multimodal agents offer a promising path to automating complex document-intensive workflows. Yet, a critical question remains: do these agents demonstrate genuine strategic reasoning, or merely stochastic trial-and-error search? To address this, we introduce MADQA, a benchmark of 2,250 human-authored questions over a





# Leaderboard

<https://huggingface.co/spaces/Snowflake/MADQA-Leaderboard>

Spaces

Snowflake/MADQA-Leaderboard

like 13

Running

Logs

App

Files

Settings

Leaderboard

Analysis

About

Submit Results

 Leaderboard

Filter by techniques/features:  
Click to filter by tags...


Select columns to display:  
Model Type × Tags × Accuracy × Attribution × Effort ×

| Model                                                                                                                                                             | Organization | Model Type                                                                            | Accuracy<br>(LLM judge) | Attribution<br>(Page F1) | Effort<br>(Kuiper) | Tags                                                      | Analyze |
|-------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------|---------------------------------------------------------------------------------------|-------------------------|--------------------------|--------------------|-----------------------------------------------------------|---------|
|  <b>Human with Oracle Retriever</b><br>Human given gold standard evidence pages. |              |                                                                                       | 99.4 ± 0.4              | —                        | —                  | Vision and Language                                       | View    |
| <b>Gemini 3.5 Flash with Mixedbread Agentic Search</b><br>Query Mixedbread Agentic mode to get top-10 documents and pass them to the LLM.                         | Mixedbread   |  api | 93.4 ± 1.3              | 84.3 ± 1.6               | (8.8)              | Agentic Semantic Search Tool Vision and Language          | View    |
| <b>Button</b><br>Hybrid retrieval (Mixedbread + BM25 + file search tool), Gemini 3.1 pro, 3-pass agentic refinement (Generate, Verify, Fix)                       | Distyl AI    |  api | 91.7 ± 1.5              | 86.9 ± 1.5               | (12.8)             | Agentic Semantic Search Tool Vision and Language          | View    |
| <b>Gemini 3.5 Flash with Mixedbread</b><br>Query Mixedbread to get top-10 documents and pass them to the LLM.                                                     | Mixedbread   |  api | 88.9 ± 1.7              | 83.4 ± 1.7               | —                  | Conventional RAG Semantic Search Tool Vision and Language | View    |

**The best LLM agents can now match human accuracy in document intelligence tasks.**

**But they solve problems 5 times less efficiently.**

**And both humans and machines hit a performance ceiling.**



#### Correctness

Is the answer correct and actionable?



#### Attribution

Is the answer well grounded?



#### Efficiency

Is the effort the agent invested justified?



# Findings

## Simple Agentic Systems Can Outperform Strong, Static RAG, e.g.:

|          | Static RAG Service                   | Acc. | BM25 Agent Acc. | $\Delta$     |
|----------|--------------------------------------|------|-----------------|--------------|
| ● Gemini | <i>Gemini File Search</i>            | 78.6 | 82.2            | <b>+3.6</b>  |
| ● GPT    | <i>OpenAI Assistants File Search</i> | 50.0 | 67.8            | <b>+17.8</b> |

| Model / Framework                        | Accuracy                         |
|------------------------------------------|----------------------------------|
| <i>Non-Agentic Systems</i>               |                                  |
| ● Gemini 3 Pro File Search               | <b>78.6 <math>\pm</math> 2.2</b> |
| Gemini 2.5 Flash File Search             | 71.8 $\pm$ 2.4                   |
| M3DocRAG                                 | 61.6 $\pm$ 2.6                   |
| GPT-5.2 (2024-08) HEAVEN                 | 52.9 $\pm$ 2.7                   |
| ● GPT-5.2 (2025-12) File Search          | 50.0 $\pm$ 2.7                   |
| GPT-5 (2025-08) File Search              | 49.6 $\pm$ 2.7                   |
| GPT-4o (2024-08) HEAVEN                  | 48.6 $\pm$ 2.7                   |
| GPT-5 Mini (2025-08) File Search         | 48.5 $\pm$ 2.7                   |
| ColBERTv2 + Llama-3.1-8B                 | 40.2 $\pm$ 2.6                   |
| <i>Agentic Systems</i>                   |                                  |
| ● Gemini 3 Pro BM25 Agent                | <b>82.2 <math>\pm</math> 2.0</b> |
| Claude Sonnet 4.5 (2025-09) BM25 Agent   | 80.6 $\pm$ 2.1                   |
| GPT-5 (2025-08) BM25 Agent               | 77.7 $\pm$ 2.2                   |
| Gemini 3 Pro RLM                         | 73.8 $\pm$ 2.3                   |
| Claude Agent Semtools                    | 72.6 $\pm$ 2.4                   |
| Claude 4.5 Sonnet (2025-09) RLM          | 70.5 $\pm$ 2.4                   |
| Claude Haiku 4.5 (2025-10) BM25 Agent    | 68.2 $\pm$ 2.5                   |
| ● GPT-5.2 (2025-12) BM25 Agent           | 67.8 $\pm$ 2.5                   |
| GPT-5 Mini (2025-08) BM25 Agent          | 66.9 $\pm$ 2.5                   |
| GLM-4.6V BM25 Agent                      | 66.1 $\pm$ 2.5                   |
| GPT-5.2 (2025-12) RLM                    | 64.2 $\pm$ 2.6                   |
| MDocAgent                                | 63.8 $\pm$ 2.6                   |
| Qwen3-VL (235B-A22B-Thinking) BM25 Agent | 60.3 $\pm$ 2.6                   |
| GPT-4.1 (2025-04) BM25 Agent             | 60.0 $\pm$ 2.6                   |
| Gemini 2.5 Flash BM25 Agent              | 58.5 $\pm$ 2.6                   |
| GPT-5 Nano (2025-08) BM25 Agent          | 58.2 $\pm$ 2.6                   |
| Qwen3-VL (8B-Thinking) BM25 Agent        | 47.3 $\pm$ 2.7                   |
| GLM-4.6V Flash BM25 Agent                | 46.0 $\pm$ 2.7                   |
| GPT-4.1 Nano (2025-04) BM25 Agent        | 19.5 $\pm$ 2.1                   |
| <i>Human Performance</i>                 |                                  |
| Human Oracle Retriever                   | <b>99.4 <math>\pm</math> 0.4</b> |
| Human BM25 Agent                         | <b>82.2 <math>\pm</math> 2.0</b> |

# Adaptive compute: effort calibration?



## Section 02

# Inner Loop

EVIDENCE

Can I trust it?





# Can I trust what I found?

## Provenance Crop-Based Agentic Reflection

### BASELINE PREDICTION

2, 5, 6, 9, 11, 12, 14, 15, 18, 19, 20 ✗

### IDP-AGENT PREDICTION

2, 5, 6, 8, 9, 11, 12, 14, 15, 18, 19 ✓

"The corrected list matches the 'Total' (合計) count of 11 provided in the image by resolving circled marks."

Agent resolved complex handwriting

First-pass  
extraction

Provenance /  
location

Crop

Inspect and  
verify

Find smallest context that can test the claim.  
Reflection turns evidence into automation (or trust).

| 施術内容    |                                                                                                                                                                                                        | 金額   | 施術内容          |            | 金額    |
|---------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------|---------------|------------|-------|
| 初検料     | (時間外・深夜・休日) 円                                                                                                                                                                                          |      | 電             | ① 550 x 11 | 6050  |
| 再検料     | 円 x 回                                                                                                                                                                                                  |      | 電             | ② 550 x 11 | 6050  |
| 指導管理料   | 680 円 x 4 回                                                                                                                                                                                            | 2720 | 電             | ③ 550 x 11 | 6050  |
| 往復料     | 距離 (片道) km<br>円 x 回                                                                                                                                                                                    |      | 料             | ④          |       |
|         |                                                                                                                                                                                                        |      | 小計            |            | 18150 |
|         |                                                                                                                                                                                                        |      | あ             | ① 95 x 11  | 1045  |
| 小計      | 2720                                                                                                                                                                                                   |      | ん             | ② 95 x 11  | 1045  |
| ①       |                                                                                                                                                                                                        |      | 法             | ③ 95 x 11  | 1045  |
| ②       |                                                                                                                                                                                                        |      | 料             | ④          |       |
| ③       |                                                                                                                                                                                                        |      | 小計            |            | 3135  |
| ④       |                                                                                                                                                                                                        |      | そ             |            |       |
| 特別材料料   | 円 x 部位<br>円 x 箇所                                                                                                                                                                                       |      | の             |            |       |
|         |                                                                                                                                                                                                        |      | 他             |            |       |
| 小計      |                                                                                                                                                                                                        |      | 施術証明書・施術費明細書料 |            | 5000  |
| ①       | 615 x 11                                                                                                                                                                                               | 6765 | 合 計           |            | 55300 |
| ②       | 615 x 11                                                                                                                                                                                               | 6765 | 社会保険への請求額     |            | 55300 |
| ③       | 615 x 11                                                                                                                                                                                               | 6765 | 患者負担 0 %      |            |       |
| ④       |                                                                                                                                                                                                        |      | 一部負担金         |            |       |
| 包等交換料   | 円                                                                                                                                                                                                      |      | 給付対象外         |            | 5000  |
| 小計      | 26295                                                                                                                                                                                                  |      | 計             |            |       |
| 請求 別    | 施術料 ¥ 5000 を                                                                                                                                                                                           |      | に請求中          |            |       |
| 通院日     | 通院の場合は必ず通院日に○印をつけて下さい。(往復は△印を)                                                                                                                                                                         |      |               |            |       |
| 11月     | 1: 0 2: 3 3: 4 4: 5 5: 6 6: 7 7: 8 8: 9 9: 10 10: 11 11: 12 12: 13 13: 14 14: 15 15: 16 16: 17 17: 18 18: 19 19: 20 20: 21 21: 22 22: 23 23: 24 24: 25 25: 26 26: 27 27: 28 28: 29 29: 30 30: 31       |      | 11日           |            |       |
| 12月     | 1: 1 2: 2 3: 3 4: 4 5: 5 6: 6 7: 7 8: 8 9: 9 10: 10 11: 11 12: 12 13: 13 14: 14 15: 15 16: 16 17: 17 18: 18 19: 19 20: 20 21: 21 22: 22 23: 23 24: 24 25: 25 26: 26 27: 27 28: 28 29: 29 30: 30 31: 31 |      | 12日           |            |       |
| 1月      | 1: 1 2: 2 3: 3 4: 4 5: 5 6: 6 7: 7 8: 8 9: 9 10: 10 11: 11 12: 12 13: 13 14: 14 15: 15 16: 16 17: 17 18: 18 19: 19 20: 20 21: 21 22: 22 23: 23 24: 24 25: 25 26: 26 27: 27 28: 28 29: 29 30: 30 31: 31 |      | 1日            |            |       |
| 取 引     | 銀行 口座                                                                                                                                                                                                  |      | 上記の通り記載いたします。 |            |       |
| 所 在 地   | 和歌山県和歌山市友田町2-2-2                                                                                                                                                                                       |      | 6年11月23日      |            |       |
| 電 話 番 号 | 025-1234-5678                                                                                                                                                                                          |      |               |            |       |
| 名 称     |                                                                                                                                                                                                        |      |               |            |       |



## Section 02

# Inner Loop

**ACTION**

**What should I do?**

# What should I do?

## ROUTING ENGINE

### Dynamic Compute Allocation

Once evidence is gathered and verified, the system must decide what should happen next and what entity should perform it. The agent routes the action to the most appropriate substrate.

## 01 Model Classes

LLMs, VLMs, and specialized models selected dynamically for reasoning and unstructured processing.

---

## 02 Deterministic Workflows

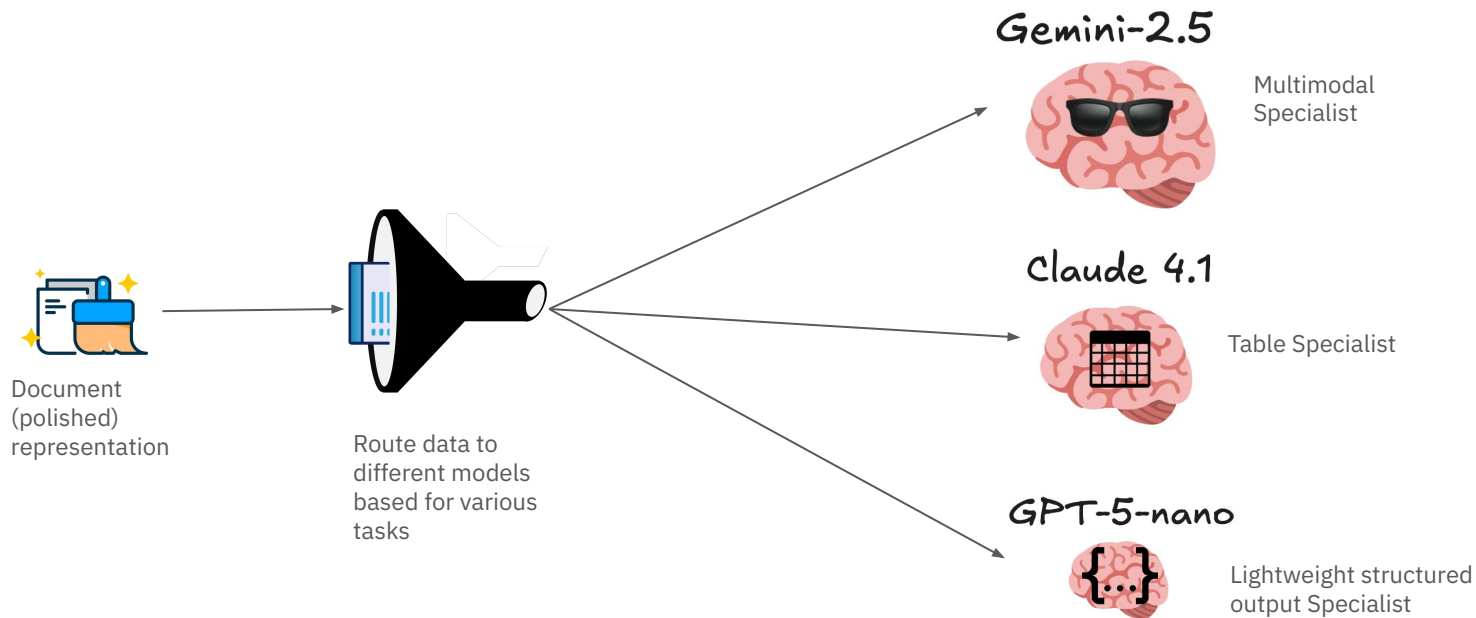
Parsers and predefined logic paths utilized when execution must remain strictly bounded and reliable.

---

## 03 Code & Human Execution

Python playbooks, direct SQL queries, or escalation to human experts when high-fidelity execution is required.

# Example 1: Routing according to complexity/cost-effectiveness



# Example 2: structured data reconciliation

## Same domain, two carrier schemas

Base Bordereau.xlsx vs. Axis Reinsurance.xlsx — identical business data, incompatible column names and sheet names.

■ claim identifier ■ gross loss / claimed amount ■ paid amount ■ outstanding / reserve ■ claim status ■ geography

Base Bordereau.xlsx

Claims

|   | Claim ID | Policy Num. | Status    | Loss Amount | Paid Amount | Currency | Outstandin.. | Country |
|---|----------|-------------|-----------|-------------|-------------|----------|--------------|---------|
| 1 | CLM10000 | POL2051     | In Review | 60,303.91   | 74,349.33   | GBP      | -14,045.42   | UK      |
| 2 | CLM10001 | POL2061     | In Review | 82,853.08   | 44,842.96   | USD      | 38,010.12    | Japan   |
| 3 | CLM10002 | POL2057     | Open      | 95,948.40   | 35,078.23   | EUR      | 60,870.17    | Japan   |
| 4 | CLM10003 | POL2051     | Open      | 34,909.90   | 40,786.14   | GBP      | -5,876.24    | Japan   |

1,000 rows

15 columns · header row 0 · claim id = Claim ID

Axis Reinsurance.xlsx

Claims Register

|   | Loss ID  | Policy Ref | Claim Stat.. | Claimed Am.. | Amount Paid | Currency C.. | Remaining .. | Territory |
|---|----------|------------|--------------|--------------|-------------|--------------|--------------|-----------|
| 1 | CLM10000 | POL2051    | Closed       | 59,097.27    | 24,048.58   | JPY          | 35,048.69    | France    |
| 2 | CLM10001 | POL2092    | Closed       | 74,798.51    | 32,732.21   | USD          | 42,066.30    | Japan     |
| 3 | CLM10002 | POL2014    | Open         | 43,734.30    | 8,344.03    | JPY          | 35,390.27    | France    |
| 4 | CLM10003 | POL2071    | In Review    | 13,630.45    | 84,357.73   | GBP          | -70,727.28   | USA       |

1,000 rows

15 columns · header row 0 · claim id = Loss ID

DIFFERENT NAMES

agents-excel /Users/jv/code/benchmarks/agents-excel/demos/05\_bordereaux\_multi\_car Connecting...

connecting...

CHAT

agent

Connect a work directory to start.

EXECUTION TRACE

Ask about the files... Send

### Example 3: boundary condition of human escalation

|                                     |                                                 |                             |                |                       |
|-------------------------------------|-------------------------------------------------|-----------------------------|----------------|-----------------------|
| Any weight gain in past few months? | <input checked="" type="checkbox"/> Yes         | <input type="checkbox"/> No | If yes: Amount | <del>due to not</del> |
| Reason:                             | <del>due to not eating</del> <del>cooking</del> |                             |                |                       |

# ExtractBench

## Input: JSON Schema

```
"county": str
"reason_pressure": bool
"api_no": str|null
... 159 fields
```

## Input: the document

## Output: JSON + evidence

```
"county": "Atascosa"
  ✓ value + word-level box
"reason_pressure": false
  ✓ unchecked box, so false
"api_no": null
  ✓ the field is blank, so null
... all 159 fields scored
```

## Scored on

**Value F1**  
array rows match in any order

**2 levels Grounding F1**  
word box + page grounding


**Per-challenge scores**  
sliced by 22 challenge tags

**Cost per page**  
0.2¢ to 34¢ across systems


W-14 disposal permit  
dense document, scanned and handwritten, energy domain

370 documents, 4,869 pages, 67 document types, 8 domains, 22 challenge tags, 14 extraction systems


### Task challenges




### Perception challenges



### Table structure





### Document length




documents, % of 370

### Business domain



Section 03

# Outer Loop

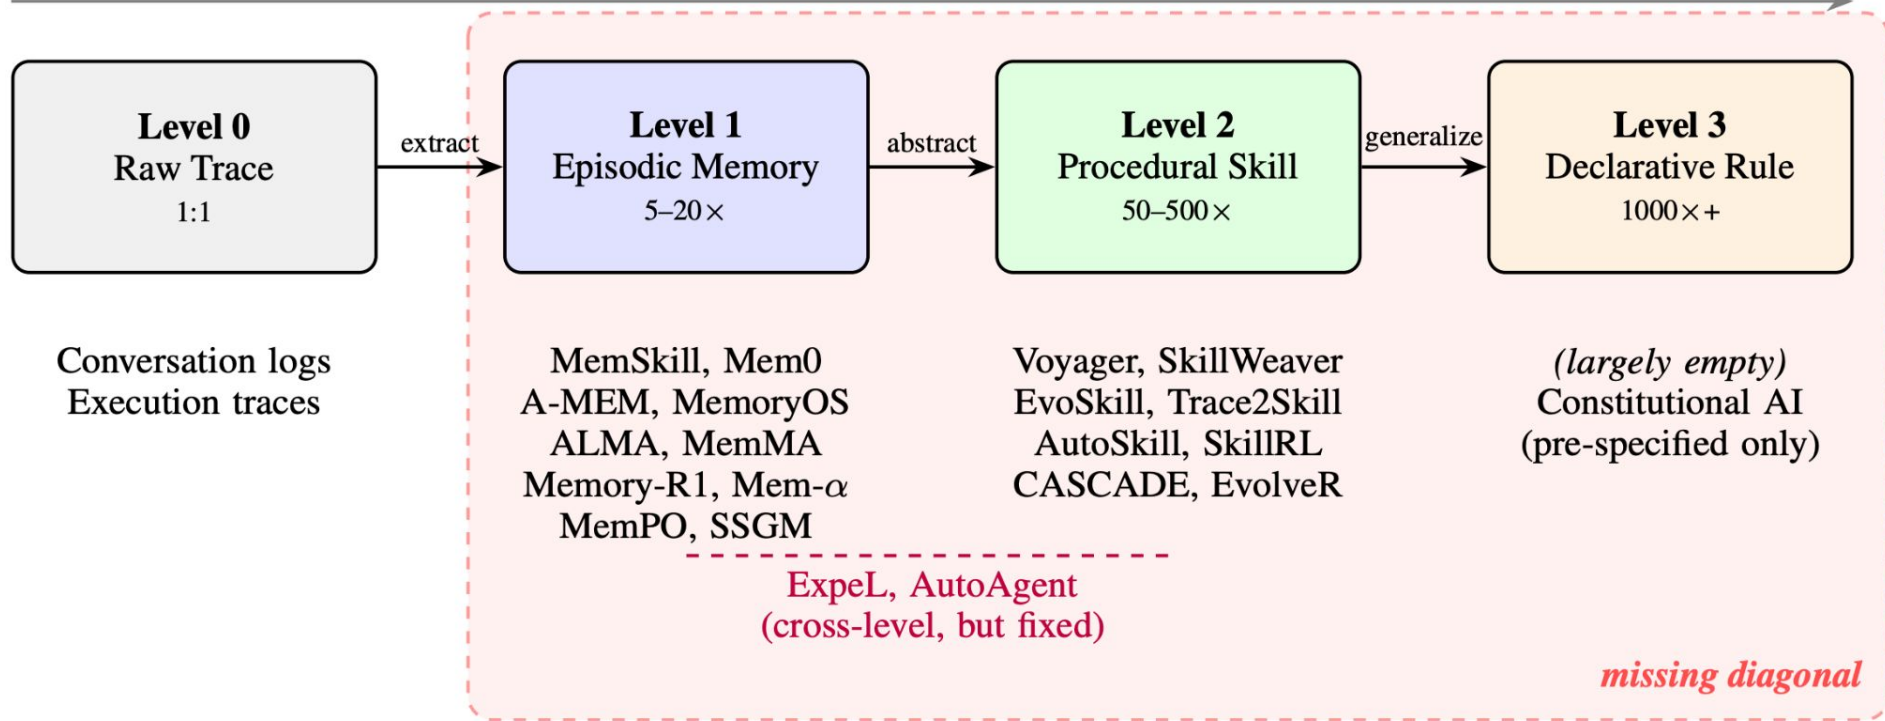




# Experience compression spectrum (Zhang et al. 2026)

← Specificity

Generalizability →



# Don't waste experience

SUCCESS

FAILURE

HUMAN CORRECTION

ESCALATION

TOOL TRACE

PRODUCTION DRIFT

## GUIDING PRINCIPLE

**Missing competence**

→ scaffold it.

**Stable competence**

→ compile it.

**Residual uncertainty**

→ keep it adaptive.

# Walkthrough example (success)

*Claimant\_would\_benefit\_from\_equipment\_description*

Based on your observation and/or clinical judgment, would the claimant benefit from using any equipment? ☒ Yes ☐ No

If yes, what is your recommendation? (check all that apply)

☐ Cane ☒ Walker ☐ Wheelchair ☐ Other: \_\_\_\_\_

Please explain: Client has obvious balance issues & fall potential

**Motivation:** "I have corrected the extracted value to better match the handwriting in the image. Specifically, the symbol between 'issues' and 'fall' is a medical shorthand for 'with' (c with a bar over it), and 'potential' was spelled out on the line below "

handwriting on evidence page in medical form → provenance reflection

# Walkthrough example (failure)

Box too tight

If claimant is able to ambulate, describe how the claimant was able to walk around the room and other areas of the residence: She has decreased balance when turning. She walks in a slow gait

→

☐ Limited range of motion ☐ Lack of strength ☐ Balance issue ☐ Cognitive impairment  
☐ Other: \_\_\_\_\_

If claimant is able to ambulate, describe how the claimant was able to walk around the room and other areas of the residence: She has decreased balance when turning. She walks in a slow gait

Balance is sufficient other than during turns. May benefit from standby assistance to avoid falls.  
Based on your observation, does the claimant experience any of the following: Two recent falls.

☒ Leaning Side-ways ☒ Shuffling Gait ☒ Staggering ☒ Unsteadiness/Imbalance ☐

handwriting on evidence page in medical form → provenance region\*150%

# Walkthrough example (failure)

| Medical Condition | Onset Date             |
|-------------------|------------------------|
| Diabetes          | 1/1/2019<br>MM DD YYYY |
| Vit B12 deficient | 1/1/2020<br>MM DD YYYY |
| Balance Issues    | 1/1/2022<br>MM DD YYYY |
| Vit D3 deficient  | 1/1/2022<br>MM DD YYYY |
| Fall potential    | 1/1/2022<br>MM DD YYYY |
| Wt loss           | 1/1/2023<br>MM DD YYYY |
| Dementia          | Unknown/<br>MM DD YYYY |
| Hypertension      | Unknown/<br>MM DD YYYY |
|                   | MM DD YYYY             |
|                   | MM DD YYYY             |
|                   | MM DD YYYY             |
|                   | MM DD YYYY             |

Claimant-Reported: Height 5 ft 2 in. Weight 102 lbs. 88 lbs

Any weight gain in past few months? ☒ Yes ☐ No If yes: Amount 8 pounds

Reason: due to not eating due to not cooking

Any weight loss in past few months? ☒ Yes ☐ No If yes: Amount 8 pounds

Reason: due to lack of interest in eating

# Walkthrough Example: (Parametric Updates)

**Train Vision-Language Models to generate better search queries through trajectory learning.**

💡 Use the **normalized rank of GT doc/pages** as a reward signal to train the model to generate more effective search queries.

Question: Search for documents about apache helicopter specification

Output: <think>

Okay, the user is asking for documents about Apache helicopter specifications. Hmm, this seems straightforward but I need to think carefully. They're probably a researcher, military enthusiast, or maybe someone in aviation industry. Could also be a student working on a project.

First thought: They didn't specify which Apache variant (UH-60 vs AH-64). That's crucial because specs differ wildly. The AH-64 Apache attack helicopter has very different specs than the UH-60 Black Hawk transport. Should clarify that immediately to avoid misinformation.

(...)

\*preparing response\*

Will start with variant clarification, then pivot to search strategies. The "search terms" part is most actionable - they'll probably copy-paste those. Should include both exact phrases ("AH-64 Block III") and broad terms ("Apache specifications").

[Think]

Okay, the user is asking for a search on Apache helicopter specifications. (...)

Apache specs usually cover things like speed, range, armament, and capabilities. I recall the AH-64E is the latest variant, but I should confirm if they want all models or just current ones. Since they didn't specify, I'll cover the main variants. (...)

Final thought: Keep it precise but offer flexibility. They might actually want maintenance manuals or flight manuals - but since they said "specifications", I'll focus on performance and capabilities first.

[Search Query]

**Apache helicopter specifications AH-64E technical data armament speed range sensor systems engine details**

# Start benchmarking failure modes.

---

## REPRESENTATION

What does faithful mean for an AI-native document substrate?

---

## SEARCH

Can trajectories teach calibrated stopping and human-agent routing?

---

## EVIDENCE

Can evidence-grounded signals outperform scalar confidence on hard fields?

---

## EXECUTION

Can the language-versus-code boundary be learned and audited?

## METHOD

**Enumerate the failure taxonomy first; synthesize one minimal case per failure.**

# We don't eliminate uncertainty.

We aim for systems that can reduce it, check it, and learn from what happens over time.



Reduce uncertainty. Don't waste experience.



# References

## Benchmarks & Datasets

- **DUDE (still not solved!)**
- **MADQA** ([arXiv:2603.12180](#))
- **Parsebench & Extractbench** [[LlamaIndex](#)]
- **Micro1 LongExtract-50** [[Reducto](#)]
- **RealDocBench** [[Extend](#)]
- **ConfBench** ([arXiv:2608.01792](#))
- **OfficeQA Pro v2** [[DataBricks](#)]
- **GDPVal** [[OpenAI](#)]
- **GPD.pdf** [[Surge AI](#)]
- **EnterpriseOps-Gym** [[ServiceNow](#)]

## Agent Architectures & Memory

- **Experience Compression Spectrum: Unifying Memory, Skills, and Rules in LLM Agents** ([arXiv:2604.15877](#))
- **SkillsSmith: Learning to Compose Parametric Skills & Knowledge** ([arXiv:2607.27497](#))
- **Compiled AI: Deterministic code generation for workflow automation** ([Trooskens et al., 2026](#))
- **Inference Time Feedback for Agents** ([arXiv:2604.27233](#))
- **SIA: Self Improving AI with Harness & Weight Updates** ([arXiv:2602.07755](#))
- **Mem- $\pi$ : Adaptive Memory through Learning When and What to Generate** ([arXiv:2605.21463](#))
- **Learning to Continually Learn via Meta-learning Agentic Memory Designs** ([arXiv:2605.27276](#))